Belief-State Control for Diplomatic Negotiation with a Lower-Cost Language Model

Anonymous ACL submission

Abstract

Can explicit strategic control improve negotiation with a lower-cost language model? We study a simulation grounded in 12 real diplomatic documents, holding GPT-5.6 Luna's generated utterances and their interpretations fixed across controllers. On 64 matched conditions, belief-state planning achieves net utility 0.533, compared with 0.514 for direct action selection by GPT-5.6 Sol with medium reasoning. Planner-distilled PPO reaches 0.475 without further model calls for strategy selection. Separate interventions attribute most of the learned policy's gain to signals and updated beliefs. Dialogue analysis connects these mechanisms to conditional offers and commitment wording. The results identify the benefits of explicit control under declared negotiation rules.

1 Introduction

Large language model (LLM) agents create opportunities to negotiate on behalf of users (Yao et al., 2023; Bianchi et al., 2024). They must pursue user goals while responding to parties with different interests and incomplete information. Existing benchmarks show that strategic choices affect outcomes (Abdelnabi et al., 2024). Repeated decisions also make inference cost relevant (Kapoor et al., 2024), motivating systems that assign specialized roles to lower-cost models (Belcak et al., 2025).

Prior work separates strategy from language generation (He et al., 2018), combines dialogue with planning (Meta Fundamental AI Research Diplomacy Team (FAIR), 2022), and improves smaller negotiators through reinforcement learning and distillation (Hua et al., 2026; Tabacof and Joglekar, 2026). Building on these approaches, we ask: with language realization held fixed, how does explicit strategic control affect negotiation utility and the cost of choosing actions?

We construct a diplomatic negotiation simulation from 12 real documents across six event groups. The documents provide context and examples of reciprocal conditions; utilities, hidden opponent types, and responses are synthetic. GPT-5.6 Luna (Luna) generates and separately interprets 396 fixed utterances. We compare Random, Heuristic, Bandit, PPO, direct model decisions, an exact public-belief planner, and a policy distilled from that planner (Figure 1). The planner knows the simulation rules but never observes the hidden type.

On 64 matched conditions, the planner achieves higher utility than direct action selection by GPT-5.6 Sol (Sol) with medium reasoning. Distilled PPO improves on base PPO and acts without further model calls for strategy selection. We then use interventions and matched dialogue traces to explain these results. This analysis distinguishes questions that reveal information from signals that change acceptance incentives, and examines whether actions preserve the conditions and commitments expressed in dialogue. Our contribution is a controlled comparison linking negotiation outcomes, strategic computation, and discourse meaning.

2 Related Work

Negotiation agents. Lewis et al. (2017) learn from human bargaining dialogues and improve decisions through simulated continuations. Recent benchmarks cover resource allocation, trading, and multiple stakeholders (Bianchi et al., 2024; Abdelnabi et al., 2024). CICERO combines dialogue and strategic reasoning in the board game Diplomacy (Meta Fundamental AI Research Diplomacy Team (FAIR), 2022). Our simulation uses real diplomatic documents to frame reciprocal measures and commitments.

Strategic control. He et al. (2018) separate dialogue acts from linguistic realization. GDP-Zero uses LLM simulations for tree search (Yu et al., 2023); PPDPP learns a separate policy planner

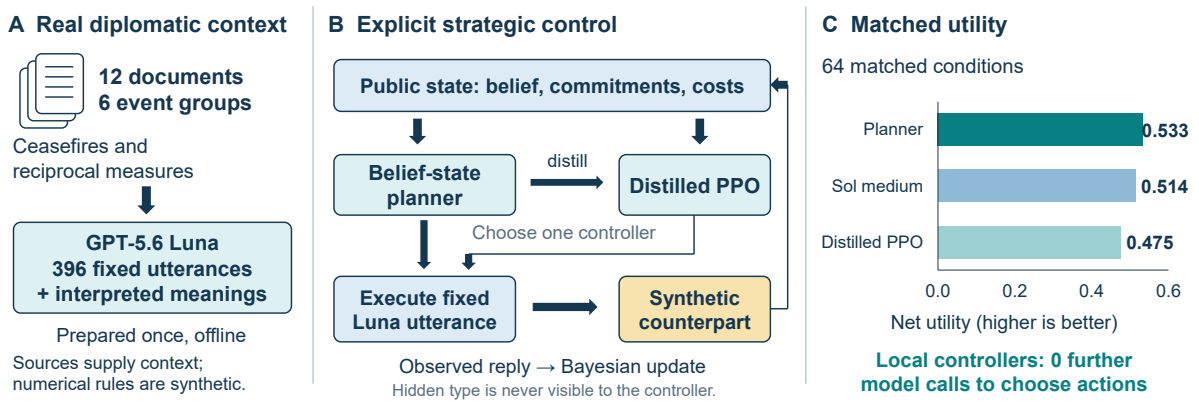


Figure 1: Overview. (A) Real diplomatic sources provide context for a fixed set of Luna utterances and interpreted meanings. (B) A public-belief planner or its distilled PPO policy selects actions; observed responses update the public state. Distillation transfers planner supervision to the policy, while language stays fixed. (C) Matched utility with the same language execution. Zero further model calls refers only to local strategy selection.

(Deng et al., 2024); and Dual-Mind combines strategy planning with expression refinement (Liu et al., 2025). We build on this separation by holding utterances and interpretations fixed when comparing controllers, then analyzing their information use and commitments.

Efficient agents. Kapoor et al. (2024) advocate evaluating performance alongside cost, while Belcak et al. (2025) argue for specialized small-model roles. SocialRL combines reinforcement learning and teacher distillation for negotiation (Hua et al., 2026). Tabacof and Joglekar (2026) study reinforcement learning across seller model sizes. We keep language model weights fixed and train a separate controller, measuring strategy-selection calls and examining the mechanisms behind utility gains.

3 Method

Data and task. We use 12 document contexts from six diplomatic event groups, including CMFA PressCon (Mochtak and Turcsanyi, 2021) and Security Council discourse. Training covers the Korean Peninsula, South Sudan, and Kashmir (six contexts); development uses Gaza (two); testing uses Syria and Iran (four). Passages about ceasefires, reciprocal measures, and commitments guide synthetic bilateral negotiations. One party offers to maintain its measure for $x \in \{1, 2, 3\}$ weeks if the counterpart maintains its measure for $y \in \{1, 2\}$ weeks. Durations, payoffs, types, and responses are researcher-defined.

State, actions, and reward. The agent has at most five decisions and cannot observe the counterpart's type, $\tau \in \{S, H\}$. It observes the offer, acceptance, commitments, costs, and belief $b_t = P(\tau = H)$. Eleven actions cover ambiguity, signaling, a private question, six conditional offers, ratification, and withdrawal. Questions provide noisy type cues; offer responses also reveal information. Ratification requires observed acceptance and own utility at least as high as the outside option.

Specifically, an agreement gives $u_A = 0.5y - 0.18x$, and the counterpart accepts when $u_B = 0.35x - c_\tau y + \delta z \geq 0$, where z records a signal and c_τ determines the cost of reciprocal measures. A signal therefore changes an incentive; a question changes information. For response likelihoods L_H and L_S, the public belief becomes

$$b_{t+1} = \frac{b_t L_H}{b_t L_H + (1 - b_t) L_S}. \quad (1)$$

Equal response likelihoods leave the belief unchanged. Step rewards are

$$r_t = d_t u_{\text{end}} - 0.01 - c_t^{\text{act}} - 0.005 e_{t+1}^2, \quad (2)$$

where d_t marks termination, e counts rejected offers, and c_t^{act} sums question, signal, rejection, reversal, and invalid-action charges. Terminal utility is u_A at agreement or 0.12 at withdrawal or expiration, paid once. We report $\sum_t r_t$. Appendix A gives all costs and transition rules.

Shared language execution. GPT-5.6 Luna generates three utterance styles for each context–action pair, yielding 396 utterances. A separate Luna call interprets each utterance's function, reciprocal terms, and stated cost. These interpreted meanings

determine transitions, and mismatches with the intended action are retained. Every controller uses this same fixed set. Counterpart replies use fixed templates. This separates action selection from language realization while keeping the consequences of interpreted utterances in the environment.

Strategic controllers. We compare Random, a development-selected signal-first Heuristic, an immediate-reward contextual Bandit, and PPO (Schulman et al., 2017). An exact belief-state planner maximizes expected net utility using the declared response rules and interpreted utterances, without access to the true type. To obtain a learned controller, we distill its action values and then refine the policy with PPO. Teacher targets come only from training contexts and configurations. The distilled policy acts without invoking the planner. We also test direct action selection by Luna with reasoning set to none, and GPT-5.6 Sol with none or `medium`. Direct models receive the public state and numerical rules.

The planner evaluates future response branches by backward recursion over the remaining decisions. Distillation uses 43,440 public states from planner trajectories and exploratory trajectories. The policy learns a soft distribution over the eleven action values, with temperature 0.03, together with a value estimate. After 2,400 supervised updates, PPO refines the same network through interaction. At deployment, it receives 31 public features and chooses its highest-scoring action. Thus, planning supplies offline supervision, while the learned policy replaces repeated planning with a local network evaluation.

4 Experimental Setup

The initial evaluation crosses four test contexts with 64 numerical configurations and three styles. A matched comparison uses the first 16 configurations and one style, giving 64 context–configuration conditions. A follow-up uses 64 new configurations with the same test contexts and three styles. We compute expectations over hidden types and question responses, then average within event groups and across groups. Learned results average three seeds. Bandit and base PPO each receive 102,400 training steps per seed; distilled PPO receives additional teacher supervision. Final checkpoints are used, with method selection on development data. Appendices A and C provide settings and diagnostic analyses.

Controller	Utility	Agree.	Calls
A. Initial: 256 conditions, three styles			
Random	-0.146 ± 0.001	11.0	0
Bandit	0.118 ± 0.001	23.8	0
Heuristic	0.419	94.8	0
PPO	0.358 ± 0.018	73.0	0
Planner	0.488	93.3	0
B. Matched: 64 conditions, one style			
Luna none	-0.368	24.3	3.949
Sol none	-0.367	26.3	4.773
Sol medium	0.514	100.0	3.429
Heuristic	0.470	100.0	0
PPO	0.378 ± 0.012	81.9	0
Distilled PPO	0.475 ± 0.016	98.3	0
Planner	0.533	100.0	0

Table 1: Policy comparison with fixed Luna utterances. Agreement is a percentage; calls are per negotiation for strategy selection. ±: three-seed SD (evaluation seeds for Random). Panels use different conditions.

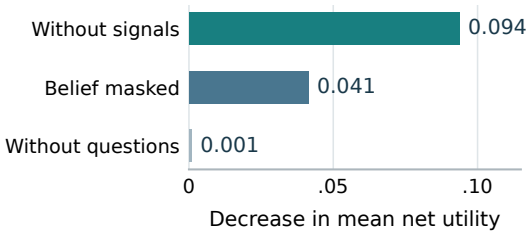


Figure 2: Evaluation-time interventions on distilled PPO. Losses are differences in three-seed mean utility on the follow-up conditions. Full values and seed variation appear in Table 3.

5 Results

Table 1 separates the initial and matched evaluations. In the initial evaluation, PPO outperforms Random and Bandit but falls below the heuristic. In the matched comparison, the planner reaches 0.533 utility, compared with 0.514 for Sol medium. Distilled PPO reaches 0.475, improving on base PPO (0.378) and slightly exceeding the heuristic (0.470). Sol none remains near Luna none, whereas Sol medium performs substantially better. Thus, the direct-model results depend strongly on the tested reasoning setting.

Local controllers require no language model calls for action selection. Sol medium uses an expected 3.429 calls per negotiation, corresponding to 4,184.8 tokens and $0.02363 at the recorded request costs. These figures count logical requests even when experimental caching avoids a repeated call. These strategy-selection figures exclude shared utterance construction, interpretation,



Figure 3: Selected paths under the same Iran-context condition and soft counterpart. Terms $x{:}y$ denote own and reciprocal durations. Values are episode net utility; the learned path uses seed 17.

and offline training.

On the follow-up configurations, distilled PPO achieves 0.468 utility, compared with 0.364 for base PPO and 0.445 for the heuristic (Figure 2). Blocking signals reduces utility by 0.094; replacing the posterior feature with the prior reduces it by 0.041. Blocking questions reduces it by only 0.001. These are evaluation-time interventions without retraining. The learned policy benefits mainly from signals and updated beliefs in this setting. Updated beliefs can come from offer responses as well as questions. The planner's no-question comparison yields a larger utility loss of 0.008.

Discourse analysis. Matched Iran-context traces illustrate this distinction (Figure 3). Luna none repeats a rejected 1:2 offer four times and withdraws, yielding -0.80. The planner and Sol medium ask a private question, send a signal, and obtain acceptance of that same offer, yielding 0.70. The question changes the belief from 0.65 to 0; the signal raises the soft counterpart's acceptance utility from -0.05 to 0.05. The learned controller instead signals and offers 2:2, yielding 0.55. Wording also matters: in a Syria-context pair, explicit withdrawal terminates the episode, while conditional withdrawal is interpreted as invalid and incurs repeated costs. Appendix B examines these paths and distinguishes expressed readiness to ratify from ratification itself.

These paths show why action order matters. The repeated 1:2 demand is rejected by both types before signaling, so rejection provides no new type information. A question can resolve uncertainty, but it does not itself make the demand acceptable. The subsequent signal changes that condition. The learned policy reaches agreement by offering a larger own concession instead. Its successful agreement therefore does not imply that it reproduces the planner's information-gathering strategy. Comparing speech-act sequences reveals this difference even when both paths end in ratification.

6 Discussion

Our results support assigning explicit strategic decisions to a separate controller when the negotiation rules can be specified. The planner uses beliefs, action costs, and future consequences to improve decisions under the same language conditions. Distillation transfers part of this behavior to a policy that acts without search or further language model calls. Its improvement over base PPO shows the value of planner supervision in this setting. Its remaining gap to Sol medium and small advantage over the heuristic distinguish exact planning from learned control.

The interventions suggest that asking more questions is not the main source of the learned policy's improvement. Signals change acceptance incentives, while responses to offers can update beliefs without a separate question. Useful information therefore depends on what the agent can subsequently do with it and what obtaining it costs. Discourse analysis adds a complementary requirement: the chosen action must be expressed with the intended conditions and commitment strength. A conditional threat to withdraw and an actual withdrawal can lead to different transitions. Strategic evaluation should therefore connect utility, information use, and the meaning of the resulting utterance.

7 Conclusion

We studied strategic control for a lower-cost language model in a diplomatic negotiation simulation grounded in real documents. With utterances held fixed, belief-state planning achieved the highest utility in the matched comparison, and planner-distilled PPO improved on base PPO without further model calls for action selection. Interventions linked the learned policy's gains mainly to signals and updated beliefs. Dialogue traces further showed how information use and commitment wording shape negotiation paths. These findings identify explicit strategic control as a useful component of negotiation systems and show how controlled comparisons can explain its benefits.

Limitations

Real documents ground the task, but payoffs and dialogues are synthetic. Evaluation covers four contexts from two event groups; extra numerical settings are not independent historical cases. Fixed utterances limit claims about online dialogue. All 44 matched style-0 meanings equal canonical templates, so utility does not establish a generation advantage. Shared-model interpretation lacks independent human validation, and some commitment expressions are misclassified.

The planner knows the response rules and interpretation mapping; distillation adds teacher supervision and computation. These resources limit comparisons with model-free learning. Model mismatch remains untested, and the tested reasoning settings do not isolate model-size effects. Three seeds and one cached model response per state do not capture historical or model-output uncertainty. Cost comparisons cover strategy-selection calls, excluding language preparation, teacher and training costs, and local computation. Detailed assumptions and procedures appear in Appendices A–B.

Ethical Considerations

Public diplomatic statements provide context, not evidence of private preferences or predictions of state behavior. Synthetic payoffs and escalation indices should not guide real diplomatic decisions. Source identifiers and transformations are preserved, and access to the supplied corpus remains restricted to the research collaboration. Language models assisted utterance generation, interpretation, manuscript drafting, and a visual concept; the final teaser was redrawn as editable shapes using recorded results. Model interpretation checks are not human annotation or evidence of independent expert agreement.

References

Sahar Abdelnabi, Amr Gomaa, Sarath Sivaprasad, Lea Schönherr, and Mario Fritz. 2024. Cooperation, competition, and maliciousness: LLM-stakeholders interactive negotiation. In *Advances in Neural Information Processing Systems*, volume 37.

Peter Belcak, Greg Heinrich, Shizhe Diao, Yonggan Fu, Xin Dong, Saurav Muralidharan, Yingyan Celine Lin, and Pavlo Molchanov. 2025. Small language models are the future of agentic AI. *Preprint*, arXiv:2506.02153.

Federico Bianchi, Patrick John Chia, Mert Yuksekgonul, Jacopo Tagliabue, Dan Jurafsky, and James Zou. 2024. How well can LLMs negotiate? NegotiationArena platform and analysis. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 3935–3951. PMLR.

Yang Deng, Wenxuan Zhang, Wai Lam, See-Kiong Ng, and Tat-Seng Chua. 2024. Plug-and-play policy planner for large language model powered dialogue agents. In *The Twelfth International Conference on Learning Representations*.

He He, Derek Chen, Anusha Balakrishnan, and Percy Liang. 2018. Decoupling strategy and generation in negotiation dialogues. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2333–2343, Brussels, Belgium. Association for Computational Linguistics.

Wenyue Hua, Zachary Huang, Tyler Payne, Safoora Yousefi, Saleema Amershi, and Asli Celikyilmaz. 2026. From passive delegates to strategic negotiators: Reinforcing social reasoning in small language models with SocialRL. *Preprint*, arXiv:2608.13787.

Sayash Kapoor, Benedikt Stroebl, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. 2024. AI agents that matter. *Preprint*, arXiv:2407.01502.

Mike Lewis, Denis Yarats, Yann Dauphin, Devi Parikh, and Dhruv Batra. 2017. Deal or no deal? end-to-end learning of negotiation dialogues. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 2443–2453, Copenhagen, Denmark. Association for Computational Linguistics.

Yutong Liu, Lida Shi, Rui Song, and Hao Xu. 2025. A dual-mind framework for strategic and expressive negotiation agent. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 23840–23860, Vienna, Austria. Association for Computational Linguistics.

Meta Fundamental AI Research Diplomacy Team (FAIR). 2022. Human-level play in the game of Diplomacy by combining language models with strategic reasoning. *Science*, 378(6624):1067–1074.

Michal Mochtak and Richard Q. Turcsanyi. 2021. Studying Chinese foreign policy narratives: Introducing the ministry of foreign affairs press conferences corpus. *Journal of Chinese Political Science*, 26:743–761.

John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. *Preprint*, arXiv:1707.06347.

Pedro Tabacof and Sagar Joglekar. 2026. Do small language models learn to negotiate? a controlled scaling study of RL-trained sellers. *Preprint*, arXiv:2610.06204.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing reasoning and acting in language models. In *The Eleventh International Conference on Learning Representations*.

Xiao Yu, Maximillian Chen, and Zhou Yu. 2023. Prompt-based Monte-Carlo tree search for goal-oriented dialogue policy planning. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7101–7125, Singapore. Association for Computational Linguistics.

A Implementation and Reproducibility

Environment. The soft-type cost is $c_S = 0.2$. A first question yields a type cue with accuracy q; later questions give no new cue. Each turn costs 0.01, each question c_q, and each rejection c_r, plus the escalation cost in the main text. Signals incur their interpreted expenditure (intended 0.06), set $z = 1$, and do not stack. Reducing an offered duration below the recorded commitment costs 0.08. Ambiguity makes the next offer provisional without erasing prior commitments. Invalid interpretations or ratifications cost 0.10. Ratification requires observed acceptance and $u_A \geq 0.12$; withdrawal or expiration pays 0.12 before costs.

We draw 96 training, 32 development, 64 initial test, and 64 new follow-up configurations over priors, question reliability/cost, rejection cost, signal effect, and hard-type cost. Exact grids are released. The 31 public features exclude source identity and hidden type. All 17 mismatches among 396 intended/interpreted action pairs remain; these are automatic checks, not human labels.

Training and aggregation. PPO uses two 64-unit layers, learning rate 3×10^{-4}, discount .95, entropy .01, rollout 256, batch 64, and ten epochs. Bandit exploration is .2. Both use seeds 17, 29, and 43. Distillation uses Adam at .001, batches of 256, value-loss weight .5, and gradient clipping .5; 35% random exploration augments teacher trajectories. Teacher computation is additional to PPO interaction. Evaluation averages styles and configurations within event groups, then weights groups equally. SD is across seeds, not historical events. Follow-up interventions appear in Table 3. Single-seed training ablations yield .346/.327 for discount .8/.99, .366/.392 for entropy 0/.1, and .377/.371 without rejection/reversal costs.

Released detail. Source sentences and provenance are in `study_results/contexts.json`. The repository includes traces, checkpoints, reproduction commands, and the unabridged supplement in `paper/supplementary_appendix_full.tex`.

B Six Discourse Cases

These are synthetic, source-grounded episodes, not historical dialogues. Cases 1–4 share Iran context `ir1` (UNSC S/2020/644, 30 June 2020), condition 0, style 0, and the soft-type evaluation branch. Cases 5–6 share Syria context `sy1` (CMFA briefing, 1 March 2016), condition 25, and the hard-type branch. Hidden types are not controller inputs.

1. Luna none. "We will suspend the specified sanctions for 1 week, conditional upon your suspension of the specified nuclear activities for 2 weeks" is rejected and repeated three times before withdrawal. Both types reject this unsignaled offer, so belief stays .65. Repetition adds costs without information; net utility is $-.80$.

2. Planner. "Privately, we ask whether you require a larger reciprocal concession" elicits a soft response, reducing belief to zero. A signal then changes counterpart utility for the 1:2 offer from $-.05$ to .05. Question, signal, offer, and ratification yield .70: the question supplies information; the signal changes incentives.

3. Sol medium. Sol follows the same path and shared utterances, also yielding .70. Its final utterance starts "We reaffirm our readiness to ratify the currently accepted agreement". The interpreter records ratification, although readiness does not itself perform that act. This is not independent language generation.

4. Distilled PPO, seed 17. The policy signals and offers "We will suspend the specified sanctions for two weeks, provided that you suspend the specified nuclear activities for the same period". Acceptance reveals the soft type; ratification follows. Utility is .55. It avoids the question but makes a larger concession. The same readiness-to-ratify mismatch applies.

5. Explicit withdrawal. After a signal and rejected 2:2 offer, the heuristic says "We hereby withdraw from these negotiations, without proposing or ratifying any agreement". The episode ends at $-.14$. Rejection reveals infeasibility; the signal is wasted because its configured effect is zero.

6. Conditional withdrawal. With another style, the same policy instead says "If these negotiations cannot advance without compromising our principles, we will withdraw without proposing or ratifying any agreement." The interpreter treats this conditional threat as invalid. Three attempts lead to expiration and $-.47$: three invalid charges and two additional turns and escalation charges explain the .33 loss.

C Diagnostic Results

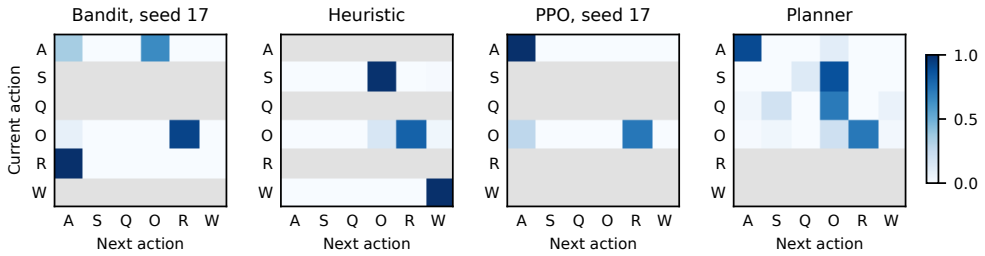


Figure 4: Row-normalized initial-test transitions. A: ambiguity; S: signal; Q: question; O: offer; R: ratify; W: withdraw. Grey: no outgoing counts. Learned panels: seed 17. Counts enumerate response paths.

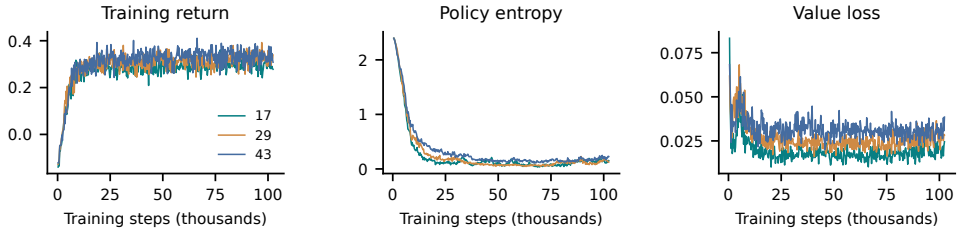


Figure 5: Base PPO training: rolling return, entropy, and value loss. Lower loss does not establish better negotiation: base PPO remains below the heuristic on the initial test.

Policy	Own	Other	Cost	Break	Turns	Max e	Mean e	Reverse	Viol.
Random	0.147	0.040	0.292	0.890	4.288	0.856	0.547	0.437	0.000
Bandit	0.140	0.082	0.022	0.762	1.773	0.030	0.017	0.000	0.000
Heuristic	0.531	0.260	0.112	0.052	3.148	0.160	0.119	0.000	0.000
PPO	0.413	0.278	0.055	0.270	2.459	0.264	0.263	0.000	0.000
Planner	0.570	0.176	0.082	0.067	3.030	0.211	0.183	0.000	0.000

Table 2: Initial-test metrics, event-group weighted. Own/Other: pre-cost utilities over all outcomes; Break: breakdown rate; e: rejection index; Reverse: reversals; Viol.: bilateral acceptance violations.

Controller / intervention	Utility	Change
PPO	$.364 \pm .016$	–
Heuristic	.445	–
Distilled only	$.446 \pm .011$	–
Distilled PPO	$.468 \pm .012$	–
Without questions	$.467 \pm .014$	$-.001$
Belief masked	$.427 \pm .013$	$-.041$
Without signals	$.374 \pm .014$	$-.094$
Planner	.500	–
Planner without questions	.492	$-.008$

Table 3: Follow-up: 256 conditions, three styles. $\pm$: three-seed SD. Changes use distilled PPO or full planner as reference. Policies are frozen; masking replaces posterior with prior.

PPO repeats ambiguity 2,260 times without gaining a type cue. Full counts and probabilities are in `transition_matrices.json`.

Changed rule	Planner	Questions	PPO
$q = 0.50$	0.527	0.00	0.410
$q = 0.70$	0.528	0.00	0.410
$q = 0.90$	0.532	1.00	0.410
$q = 1.00$	0.555	1.00	0.410
$\delta = 0.00$	0.486	1.00	0.410
$\delta = 0.10$	0.532	1.00	0.410
$\delta = 0.25$	0.596	1.00	0.410

Table 4: One-factor rule shifts: utility and expected planner question count. PPO averages three frozen seeds. Other rules stay fixed.

Base PPO uses neither questions nor signals here. Reliable information benefits the planner. Full sweeps are in `sensitivity_metrics.json`.