Belief-State Control for Diplomatic Negotiation with a Lower-Cost Language Model

Anonymous ACL submission

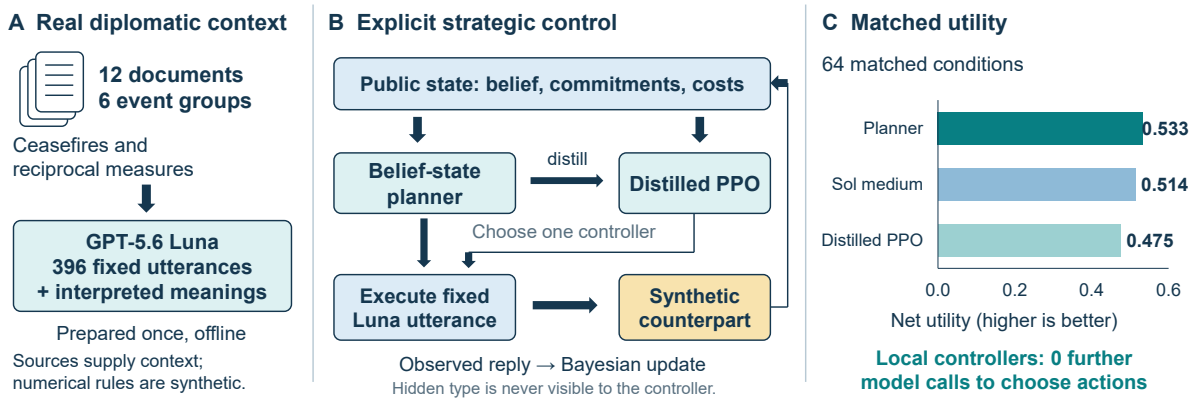


Figure 1: Real diplomatic sources ground fixed Luna utterances. A public-belief planner or its distilled PPO policy selects actions under the same language execution. Responses update public beliefs; distillation transfers planner supervision. Zero further model calls refers only to local strategy selection.

Abstract

Can explicit strategic control improve negotiation with a lower-cost language model? We study a simulation grounded in 12 real diplomatic documents, holding GPT-5.6 Luna's generated utterances and their interpretations fixed across controllers. On 64 matched conditions, belief-state planning achieves net utility 0.533, compared with 0.514 for direct action selection by GPT-5.6 Sol with medium reasoning. Planner-distilled PPO reaches 0.475 without further model calls for strategy selection. Separate interventions attribute most of the learned policy's gain to signals and updated beliefs. Dialogue analysis connects these mechanisms to conditional offers and commitment wording. The results identify the benefits of explicit control under declared negotiation rules.

1 Introduction

Large language model (LLM) agents create opportunities to negotiate on behalf of users (Yao et al., 2023; Bianchi et al., 2024; Zhan et al., 2024). They must pursue user goals while reasoning about others' interests and incomplete information (Gandhi et al., 2023; Chan et al., 2024). Benchmarks expose strategic decision challenges (Abdelnabi et al., 2024; Lin et al., 2024). Repeated decisions make inference cost relevant (Kapoor et al., 2024; Hakimov et al., 2026), motivating specialized lower-cost models (Belcak et al., 2025).

Prior work separates strategy from language generation (He et al., 2018; Deng et al., 2024), combines dialogue with planning (Meta Fundamental AI Research Diplomacy Team (FAIR), 2022; Yu et al., 2023), and improves smaller negotiators through reinforcement learning and distillation (Hua et al., 2026; Tabacof and Joglekar, 2026). We ask: with language held fixed, how does explicit control affect negotiation utility and action-selection cost?

We construct a diplomatic negotiation simulation from 12 real documents across six event groups. The documents provide context and examples of reciprocal conditions; utilities, hidden opponent types, and responses are synthetic. GPT-5.6 Luna (Luna) generates and separately interprets 396 fixed utterances. We compare Random, Heuristic, Bandit, PPO, direct model decisions, an exact public-belief planner, and a policy distilled from that planner (Figure 1). The planner knows the simulation rules but never observes the hidden type.

On 64 matched conditions, the planner outperforms direct action selection by GPT-5.6 Sol (Sol) with medium reasoning. Distilled PPO improves on base PPO without further strategy-selection model calls. Interventions and dialogue traces distinguish information-seeking questions from incentive-changing signals and test whether actions preserve expressed commitments. Our contribution links negotiation outcomes, strategic computation, and discourse meaning under controlled language conditions.

2 Related Work

Negotiation agents. Negotiation dialogue research covers bargaining, multi-party interaction, and human assistance (Zhan et al., 2024; Lin et al., 2024). Lewis et al. (2017) learn from bargaining dialogues and simulated continuations. Recent benchmarks evaluate economic outcomes and mental-state reasoning (Bianchi et al., 2024; Abdelnabi et al., 2024; Chan et al., 2024). CICERO combines dialogue and strategy in the board game Diplomacy (Meta Fundamental AI Research Diplomacy Team (FAIR), 2022); our contexts come from real diplomatic documents.

Strategic control. Separating dialogue acts from language supports explicit strategic control (He et al., 2018; Deng et al., 2024; Liu et al., 2025). Other methods use tree search, reasoning demonstrations, or diverse user simulation (Yu et al., 2023; Gandhi et al., 2023; Zhang et al., 2024). Belief-state dialogue management explicitly models uncertainty (Young et al., 2013). We compare controllers under fixed utterances and interpretations.

Efficient agents. Agent evaluation must consider performance and cost (Kapoor et al., 2024; Hakimov et al., 2026). Policy distillation and Expert Iteration transfer expert decisions into learned policies (Rusu et al., 2016; Anthony et al., 2017); recent work trains smaller negotiation models (Hua et al., 2026; Tabacof and Joglekar, 2026). We instead fix language-model weights and train a separate controller, following the motivation for specialized small-model roles (Belcak et al., 2025).

3 Method

Data and task. We use 12 document contexts from six diplomatic event groups, including CMFA PressCon (Mochtak and Turcsanyi, 2021) and Security Council discourse. Training covers the Korean Peninsula, South Sudan, and Kashmir (six contexts); development uses Gaza (two); testing uses Syria and Iran (four). Passages about ceasefires, reciprocal measures, and commitments guide synthetic bilateral negotiations. One party offers to maintain its measure for $x \in \{1, 2, 3\}$ weeks if the counterpart maintains its measure for $y \in \{1, 2\}$ weeks. Durations, payoffs, types, and responses are researcher-defined.

State, actions, and reward. The agent has at most five decisions and cannot observe the counterpart's type, $\tau \in \{S, H\}$. It observes the offer, acceptance, commitments, costs, and belief $b_t = P(\tau = H)$. Eleven actions cover ambiguity, signaling, a private question, six conditional offers, ratification, and withdrawal. Questions provide noisy type cues; offer responses also reveal information. Ratification requires observed acceptance and own utility at least as high as the outside option.

Specifically, an agreement gives $u_A = 0.5y - 0.18x$, and the counterpart accepts when $u_B = 0.35x - c_\tau y + \delta z \geq 0$, where z records a signal and c_τ determines the cost of reciprocal measures. A signal therefore changes an incentive; a question changes information. For response likelihoods L_H and L_S, the public belief becomes

$$b_{t+1} = \frac{b_t L_H}{b_t L_H + (1 - b_t) L_S}. \quad (1)$$

Equal response likelihoods leave the belief unchanged. Step rewards are

$$r_t = d_t u_{\text{end}} - 0.01 - c_t^{\text{act}} - 0.005 e_{t+1}^2, \quad (2)$$

where d_t marks termination, e counts rejected offers, and c_t^{act} sums question, signal, rejection, reversal, and invalid-action charges. Terminal utility is u_A at agreement or 0.12 at withdrawal or expiration, paid once. We report $\sum_t r_t$. Appendix A gives all costs and transition rules.

Shared language execution. GPT-5.6 Luna generates three styles per context–action pair, yielding 396 utterances. A separate Luna call interprets their function, reciprocal terms, and stated cost. These meanings determine transitions; interpretation mismatches are retained. All controllers share these utterances and templated counterpart replies, separating action selection from language realization.

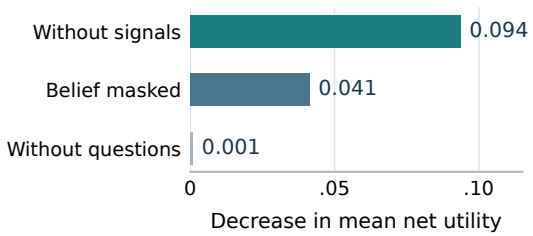


Figure 2: Evaluation-time interventions on distilled PPO. Losses are differences in three-seed mean utility on the follow-up conditions. Full values and seed variation appear in Table 3.

Controller	Utility	Agree.	Calls
A. Initial: 256 conditions, three styles			
Random	-0.146 ± 0.001	11.0	0
Bandit	0.118 ± 0.001	23.8	0
Heuristic	0.419	94.8	0
PPO	0.358 ± 0.018	73.0	0
Planner	0.488	93.3	0
B. Matched: 64 conditions, one style			
Luna none	-0.368	24.3	3.949
Sol none	-0.367	26.3	4.773
Sol medium	0.514	100.0	3.429
Heuristic	0.470	100.0	0
PPO	0.378 ± 0.012	81.9	0
Distilled PPO	0.475 ± 0.016	98.3	0
Planner	0.533	100.0	0

Table 1: Policy comparison with fixed Luna utterances. Agreement is a percentage; calls are per negotiation for strategy selection. ±: three-seed SD (evaluation seeds for Random). Panels use different conditions.

Strategic controllers. We compare Random, a development-selected signal-first Heuristic, an immediate-reward Bandit, and PPO (Schulman et al., 2017). An exact belief-state planner maximizes expected net utility under the declared rules without seeing the true type. We distill its action values from training contexts and configurations, then refine the policy with PPO. Direct Luna (reasoning `none`) and Sol (`none`/`medium`) receive the public state and numerical rules.

The planner solves response branches by backward recursion. Following policy distillation and expert-guided learning (Rusu et al., 2016; Anthony et al., 2017), we use 43,440 public states from planner and exploratory trajectories to learn soft action-value targets (temperature 0.03) and a value estimate. After 2,400 supervised updates, PPO refines the network. Deployment selects the highest-scoring action from 31 public features.

4 Experimental Setup

Initial evaluation crosses four contexts, 64 configurations, and three styles. Matched evaluation uses the first 16 configurations and one style (64 conditions); follow-up uses 64 new configurations and three styles. We take expectations over types and question responses, then average within and across event groups. Bandit and base PPO receive 102,400 steps per seed; distillation adds teacher supervision. Learned results average three seeds, using final checkpoints and development-selected methods. Appendices A and C give settings and diagnostics.

5 Results

Initially, PPO outperforms Random and Bandit but falls below the heuristic (Table 1). In the matched comparison, the planner reaches 0.533 utility versus 0.514 for Sol medium. Distilled PPO reaches 0.475 versus 0.378 for base PPO and 0.470 for the heuristic. Sol none remains near Luna none: direct-model performance depends strongly on reasoning settings.

Local controllers require no language model calls for action selection. Sol medium uses an expected 3.429 calls per negotiation, corresponding to 4,184.8 tokens and $0.02363 at the recorded request costs. These figures count logical requests even when experimental caching avoids a repeated call. These strategy-selection figures exclude shared utterance construction, interpretation, and offline training.

On the follow-up configurations, distilled PPO achieves 0.468 utility, compared with 0.364 for base PPO and 0.445 for the heuristic (Figure 2). Blocking signals reduces utility by 0.094; replacing the posterior feature with the prior reduces it by 0.041. Blocking questions reduces it by only 0.001. These are evaluation-time interventions without retraining. The learned policy benefits mainly from signals and updated beliefs in this setting. Updated beliefs can come from offer responses as well as questions. The planner's no-question comparison yields a larger utility loss of 0.008.

Discourse analysis. Matched Iran-context traces illustrate this distinction (Figure 3). Luna none repeats a rejected 1:2 offer four times and withdraws, yielding -0.80. The planner and Sol medium ask a private question, send a signal, and obtain acceptance of that same offer, yielding 0.70. The question changes the belief from 0.65 to 0; the



Figure 3: Selected paths under the same Iran-context condition and soft counterpart. Terms x:y denote own and reciprocal durations. Values are episode net utility; the learned path uses seed 17.

signal raises the soft counterpart's acceptance utility from -0.05 to 0.05. The learned controller instead signals and offers 2:2, yielding 0.55. Wording also matters: in a Syria-context pair, explicit withdrawal terminates the episode, while conditional withdrawal is interpreted as invalid and incurs repeated costs. Appendix B examines these paths and distinguishes expressed readiness to ratify from ratification itself.

Both types reject the unsignaled 1:2 demand, so repetition provides no new information. A question resolves uncertainty; the subsequent signal makes the demand acceptable. The learned policy instead concedes more. Successful agreement therefore does not establish that it reproduces the planner's information-gathering strategy.

6 Discussion

Explicit control improves decisions when negotiation rules can be specified. The planner accounts for beliefs, costs, and future responses. Distillation transfers part of this behavior to a policy without online search or model calls. Its remaining gap to Sol medium and small advantage over the heuristic distinguish exact planning from learned control.

Signals drive more of the learned gain than questions. Offer responses can also update beliefs, so information value depends on subsequent choices and acquisition costs. Discourse analysis adds a separate requirement: utterances must preserve the selected conditions and commitment strength. Evaluation should connect these meanings to utility and information use.

7 Conclusion

With Luna utterances held fixed, belief-state planning achieved the highest matched utility, and planner-distilled PPO improved on base PPO without further model calls for action selection. Signals, updated beliefs, and commitment wording explain different parts of the observed behavior. This controlled study identifies explicit strategic control as a useful component for source-grounded diplomatic negotiation simulations.

Limitations

Payoffs and dialogues are synthetic despite real-document grounding. Four test contexts cover only two event groups. Fixed utterances limit online-dialogue claims; all 44 matched style-0 meanings equal canonical templates, so utility cannot establish a generation advantage. Shared-model interpretation lacks human validation and misclassifies some commitments. Models also assisted drafting and a visual concept.

Known rules and teacher computation advantage planning and distillation over model-free learning. Model mismatch is untested; reasoning settings do not isolate model size. Three seeds and one cached response per state omit historical and output uncertainty. Reported costs exclude language preparation, teacher supervision, training, and local computation. Appendices A–B give details.

References

Sahar Abdelnabi, Amr Gomaa, Sarath Sivaprasad, Lea Schönherr, and Mario Fritz. 2024. Cooperation, competition, and maliciousness: LLM-stakeholders interactive negotiation. In *Advances in Neural Information Processing Systems*, volume 37.

Thomas Anthony, Zheng Tian, and David Barber. 2017. Thinking fast and slow with deep learning and tree search. In *Advances in Neural Information Processing Systems*, volume 30.

Peter Belcak, Greg Heinrich, Shizhe Diao, Yonggan Fu, Xin Dong, Saurav Muralidharan, Yingyan Celine Lin, and Pavlo Molchanov. 2025. Small language models are the future of agentic AI. *Preprint*, arXiv:2506.02153.

Federico Bianchi, Patrick John Chia, Mert Yuksekgonul, Jacopo Tagliabue, Dan Jurafsky, and James Zou. 2024. How well can LLMs negotiate? NegotiationArena platform and analysis. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 3935–3951. PMLR.

Chunkit Chan, Cheng Jiayang, Yauwai Yim, Zheye Deng, Wei Fan, Haoran Li, Xin Liu, Hongming Zhang, Weiqi Wang, and Yangqiu Song. 2024. NegotiationToM: A benchmark for stress-testing machine theory of mind on negotiation surrounding. In

Findings of the Association for Computational Linguistics: EMNLP 2024, pages 4211–4241, Miami, Florida, USA. Association for Computational Linguistics.

Yang Deng, Wenxuan Zhang, Wai Lam, See-Kiong Ng, and Tat-Seng Chua. 2024. Plug-and-play policy planner for large language model powered dialogue agents. In *The Twelfth International Conference on Learning Representations*.

Kanishk Gandhi, Dorsa Sadigh, and Noah D. Goodman. 2023. Strategic reasoning with language models. In *NeurIPS Workshop on Foundation Models for Decision Making*.

Sherzod Hakimov, Roland Bernard, Tim Leiber, Karl Osswald, Kristina Richert, Ruilin Yang, Raffaella Bernardi, and David Schlangen. 2026. The price of thought: A multilingual analysis of reasoning, performance, and cost of negotiation in large language models. In *Findings of the Association for Computational Linguistics: EACL 2026*, pages 529–570, Rabat, Morocco. Association for Computational Linguistics.

He He, Derek Chen, Anusha Balakrishnan, and Percy Liang. 2018. Decoupling strategy and generation in negotiation dialogues. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2333–2343, Brussels, Belgium. Association for Computational Linguistics.

Wenyue Hua, Zachary Huang, Tyler Payne, Safoora Yousefi, Saleema Amershi, and Asli Celikyilmaz. 2026. From passive delegates to strategic negotiators: Reinforcing social reasoning in small language models with SocialRL. *Preprint*, arXiv:2608.13787.

Sayash Kapoor, Benedikt Stroebl, Zachary S. Siegel, Nitya Nadgir, and Arvind Narayanan. 2024. AI agents that matter. *Preprint*, arXiv:2407.01502.

Mike Lewis, Denis Yarats, Yann Dauphin, Devi Parikh, and Dhruv Batra. 2017. Deal or no deal? end-to-end learning of negotiation dialogues. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 2443–2453, Copenhagen, Denmark. Association for Computational Linguistics.

Jessy Lin, Nicholas Tomlin, Jacob Andreas, and Jason Eisner. 2024. Decision-oriented dialogue for human-AI collaboration. *Transactions of the Association for Computational Linguistics*, 12:892–911.

Yutong Liu, Lida Shi, Rui Song, and Hao Xu. 2025. A dual-mind framework for strategic and expressive negotiation agent. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 23840–23860, Vienna, Austria. Association for Computational Linguistics.

Meta Fundamental AI Research Diplomacy Team (FAIR). 2022. Human-level play in the game of Diplomacy by combining language models with strategic reasoning. *Science*, 378(6624):1067–1074.

Michal Mochtak and Richard Q. Turcsanyi. 2021. Studying Chinese foreign policy narratives: Introducing the ministry of foreign affairs press conferences corpus. *Journal of Chinese Political Science*, 26:743–761.

Andrei A. Rusu, Sergio Gomez Colmenarejo, Caglar Gulcehre, Guillaume Desjardins, James Kirkpatrick, Razvan Pascanu, Volodymyr Mnih, Koray Kavukcuoglu, and Raia Hadsell. 2016. Policy distillation. In *International Conference on Learning Representations*.

John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. *Preprint*, arXiv:1707.06347.

Pedro Tabacof and Sagar Joglekar. 2026. Do small language models learn to negotiate? a controlled scaling study of RL-trained sellers. *Preprint*, arXiv:2610.06204.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2023. ReAct: Synergizing reasoning and acting in language models. In *The Eleventh International Conference on Learning Representations*.

Steve Young, Milica Gašić, Blaise Thomson, and Jason D. Williams. 2013. POMDP-based statistical spoken dialog systems: A review. *Proceedings of the IEEE*, 101(5):1160–1179.

Xiao Yu, Maximillian Chen, and Zhou Yu. 2023. Prompt-based Monte-Carlo tree search for goal-oriented dialogue policy planning. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7101–7125, Singapore. Association for Computational Linguistics.

Haolan Zhan, Yufei Wang, Zhuang Li, Tao Feng, Yuncheng Hua, Suraj Sharma, Lizhen Qu, Zhaleh Semnani Azad, Ingrid Zukerman, and Gholamreza Haffari. 2024. Let's negotiate! a survey of negotiation dialogue systems. In *Findings of the Association for Computational Linguistics: EACL 2024*, pages 2019–2031, St. Julian's, Malta. Association for Computational Linguistics.

Tong Zhang, Chen Huang, Yang Deng, Hongru Liang, Jia Liu, Zujie Wen, Wenqiang Lei, and Tat-Seng Chua. 2024. Strength lies in differences! improving strategy planning for non-collaborative dialogues via diversified user simulation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 424–444, Miami, Florida, USA. Association for Computational Linguistics.

A Implementation and Reproducibility

Environment. The soft-type cost is $c_S = 0.2$. A first question yields a type cue with accuracy q; later questions give no new cue. Each turn costs 0.01, each question c_q, and each rejection c_r, plus the escalation cost in the main text. Signals incur their interpreted expenditure (intended 0.06), set $z = 1$, and do not stack. Reducing an offered duration below the recorded commitment costs 0.08. Ambiguity makes the next offer provisional without erasing prior commitments. Invalid interpretations or ratifications cost 0.10. Ratification requires observed acceptance and $u_A \geq 0.12$; withdrawal or expiration pays 0.12 before costs.

We draw 96 training, 32 development, 64 initial test, and 64 new follow-up configurations over priors, question reliability/cost, rejection cost, signal effect, and hard-type cost. Exact grids are released. The 31 public features exclude source identity and hidden type. All 17 mismatches among 396 intended/interpreted action pairs remain; these are automatic checks, not human labels.

Training and aggregation. PPO uses two 64-unit layers, learning rate 3×10^{-4}, discount .95, entropy .01, rollout 256, batch 64, and ten epochs. Bandit exploration is .2. Both use seeds 17, 29, and 43. Distillation uses Adam at .001, batches of 256, value-loss weight .5, and gradient clipping .5; 35% random exploration augments teacher trajectories. Teacher computation is additional to PPO interaction. Evaluation averages styles and configurations within event groups, then weights groups equally. SD is across seeds, not historical events. Follow-up interventions appear in Table 3. Single-seed training ablations yield .346/.327 for discount .8/.99, .366/.392 for entropy 0/.1, and .377/.371 without rejection/reversal costs.

Released detail. Source sentences and provenance are in `study_results/contexts.json`. The repository includes traces, checkpoints, reproduction commands, and the unabridged supplement in `paper/supplementary_appendix_full.tex`.

B Six Discourse Cases

These are synthetic, source-grounded episodes, not historical dialogues. Cases 1–4 share Iran context `ir1` (UNSC S/2020/644, 30 June 2020), condition 0, style 0, and the soft-type evaluation branch. Cases 5–6 share Syria context `sy1` (CMFA briefing, 1 March 2016), condition 25, and the hard-type branch. Hidden types are not controller inputs.

1. Luna none. "We will suspend the specified sanctions for 1 week, conditional upon your suspension of the specified nuclear activities for 2 weeks" is rejected and repeated three times before withdrawal. Both types reject this unsignaled offer, so belief stays .65. Repetition adds costs without information; net utility is $-.80$.

2. Planner. "Privately, we ask whether you require a larger reciprocal concession" elicits a soft response, reducing belief to zero. A signal then changes counterpart utility for the 1:2 offer from $-.05$ to .05. Question, signal, offer, and ratification yield .70: the question supplies information; the signal changes incentives.

3. Sol medium. Sol follows the same path and shared utterances, also yielding .70. Its final utterance starts "We reaffirm our readiness to ratify the currently accepted agreement". The interpreter records ratification, although readiness does not itself perform that act. This is not independent language generation.

4. Distilled PPO, seed 17. The policy signals and offers "We will suspend the specified sanctions for two weeks, provided that you suspend the specified nuclear activities for the same period". Acceptance reveals the soft type; ratification follows. Utility is .55. It avoids the question but makes a larger concession. The same readiness-to-ratify mismatch applies.

5. Explicit withdrawal. After a signal and rejected 2:2 offer, the heuristic says "We hereby withdraw from these negotiations, without proposing or ratifying any agreement". The episode ends at $-.14$. Rejection reveals infeasibility; the signal is wasted because its configured effect is zero.

6. Conditional withdrawal. With another style, the same policy instead says "If these negotiations cannot advance without compromising our principles, we will withdraw without proposing or ratifying any agreement." The interpreter treats this conditional threat as invalid. Three attempts lead to expiration and $-.47$: three invalid charges and two additional turns and escalation charges explain the .33 loss.

C Diagnostic Results

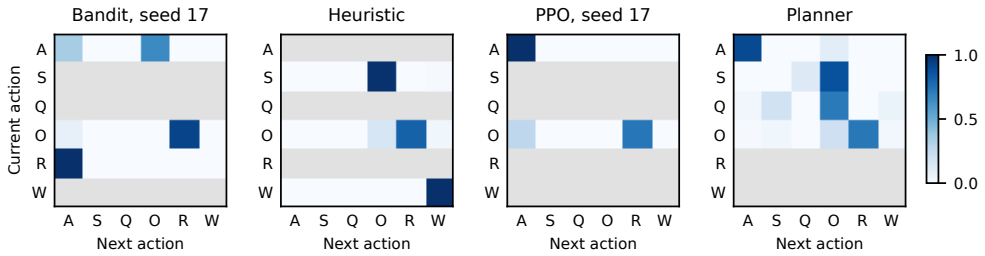


Figure 4: Row-normalized initial-test transitions. A: ambiguity; S: signal; Q: question; O: offer; R: ratify; W: withdraw. Grey: no outgoing counts. Learned panels: seed 17. Counts enumerate response paths.

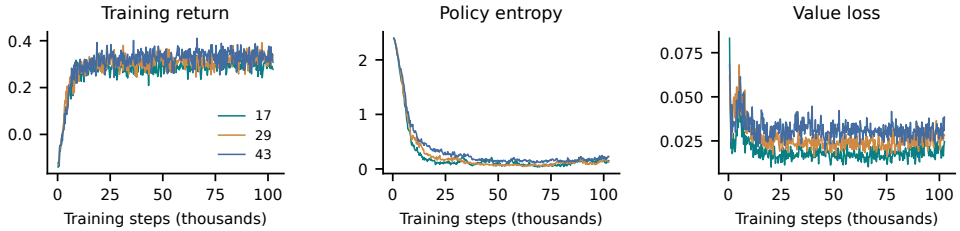


Figure 5: Base PPO training: rolling return, entropy, and value loss. Lower loss does not establish better negotiation: base PPO remains below the heuristic on the initial test.

Policy	Own	Other	Cost	Break	Turns	Max e	Mean e	Reverse	Viol.
Random	0.147	0.040	0.292	0.890	4.288	0.856	0.547	0.437	0.000
Bandit	0.140	0.082	0.022	0.762	1.773	0.030	0.017	0.000	0.000
Heuristic	0.531	0.260	0.112	0.052	3.148	0.160	0.119	0.000	0.000
PPO	0.413	0.278	0.055	0.270	2.459	0.264	0.263	0.000	0.000
Planner	0.570	0.176	0.082	0.067	3.030	0.211	0.183	0.000	0.000

Table 2: Initial-test metrics, event-group weighted. Own/Other: pre-cost utilities over all outcomes; Break: breakdown rate; e: rejection index; Reverse: reversals; Viol.: bilateral acceptance violations.

Controller / intervention	Utility	Change
PPO	$.364 \pm .016$	–
Heuristic	.445	–
Distilled only	$.446 \pm .011$	–
Distilled PPO	$.468 \pm .012$	–
Without questions	$.467 \pm .014$	$-.001$
Belief masked	$.427 \pm .013$	$-.041$
Without signals	$.374 \pm .014$	$-.094$
Planner	.500	–
Planner without questions	.492	$-.008$

Table 3: Follow-up: 256 conditions, three styles. $\pm$: three-seed SD. Changes use distilled PPO or full planner as reference. Policies are frozen; masking replaces posterior with prior.

PPO repeats ambiguity 2,260 times without gaining a type cue. Full counts and probabilities are in `transition_matrices.json`.

Changed rule	Planner	Questions	PPO
$q = 0.50$	0.527	0.00	0.410
$q = 0.70$	0.528	0.00	0.410
$q = 0.90$	0.532	1.00	0.410
$q = 1.00$	0.555	1.00	0.410
$\delta = 0.00$	0.486	1.00	0.410
$\delta = 0.10$	0.532	1.00	0.410
$\delta = 0.25$	0.596	1.00	0.410

Table 4: One-factor rule shifts: utility and expected planner question count. PPO averages three frozen seeds. Other rules stay fixed.

Base PPO uses neither questions nor signals here. Reliable information benefits the planner. Full sweeps are in `sensitivity_metrics.json`.